VoxClinBench: A Cross-Lingual, Cross-Disease Clinical Voice Biomarker Benchmark

Anonymous Author(s)
Affiliation
Address
email

Abstract

Voice-based general-practitioner (GP) diagnosis — a clinician or AI that listens once and infers a broad range of conditions — is an aspirational clinical AI target. Two first-principles arguments ground any benchmark aimed at it. (i) Physical (laryngeal/motor) and psychiatric voice conditions foreground different acoustic substrates (articulatory/phonatory kinematics vs prosodic/spectral affect) and must be evaluated in separate families, not as a unified joint prediction. (ii) Cross-lingual evaluation is a load-bearing dimension of any voice biomarker claim: a model that performs well across languages is stronger than one that only works in its training language. Every cross-corpus row in VoxClinBench tests this directly. No existing benchmark operationalises both commitments across multiple disease families and multiple languages under a common protocol. Single-disease challenges (ADReSSo [Pan et al., 2021], AVEC [Ringeval et al., 2019], FEMH [Koudounas et al., 2024], CogBench [Feng et al., 2025], SLP-Helm [Patel et al., 2025]) cover narrow niches; prior multi-disease baselines on B2AI (MARVEL [VoxClinBench Authors, 2026]) report single-seed, single-model evaluations that hide model–disease disagreement and can produce phantom leaderboard claims on small-N psychiatric tasks. We introduce **VoxClinBench**: 18 subject-disjoint tasks (16 within-corpus + 2 cross-lingual LODO) across 12 conditions and **4 languages (English, Spanish, German, Mandarin)** over **five** public or credentialed corpora (B2AI v3.0, NeuroVoz, SVD, E-DAIC, MODMA; E-DAIC is the AVEC 2019 superset of the original DAIC-WOZ [Gratch et al., 2014] cohort plus 86 new interviews; MODMA is retained as a *cross-lingual validation target only* (n=52 Chinese, small-sample caveat; not used as a training source for depression LODO in v0.3)), with a 5-seed bootstrap + DeLong [DeLong et al., 1988, Sun and Xu, 2014] + Holm–Bonferroni protocol stack and a panel of reference models (frozen SSL across multiple backbones and layers, LP-FT, handcrafted AVQI, and an 8-modality CNN–Transformer). The benchmark does not enforce a single aggregated ranking: it lists operating points for every reference model in parallel so that readers can see where languages fail to transfer, where different methods win different conditions, and where small-sample protocols mask effect sizes. Under the 5-seed patient-level bootstrap, our 8-modality reference baseline posts **macro-physical-8** $= 0.834 \pm 0.008$ (n=8 phonation-organ diseases) and **macro-psychiatric-4** $= 0.684 \pm 0.061$ (n=4 psychiatric conditions) on B2AI Tier-2, with combined macro-12 $= 0.784 \pm 0.020$ (95% CIs in supp §2); these two family-level numbers are the primary headline, reported in parallel rather than collapsed, because physical and psychiatric voice disorders are not a single skill. Representative findings that the shared protocol renders visible include a multilingual-SSL directional advantage on non-English cohorts ($+0.02$–0.04 AUROC on NeuroVoz and SVD); a disease-dependent cross-lingual split on physical tasks — B2AI$\rightarrow$NeuroVoz Parkinson's transfers zero-shot at 0.850 AUROC while B2AI$\leftrightarrow$SVD VF-paralysis collapses to ≈ 0.46 in both

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

directions; a paired GL-vs-raw causal probe on E-DAIC (N=275, Δ=+0.018) and SVD (N=513, Δ=+0.032, p<0.001) showing Griffin–Lim reconstruction does not degrade mean-pooled SSL probe AUROC on these two corpora; and a full-FT catastrophic-forgetting collapse (0.45–0.56) on small-N Griffin–Lim-reconstructed B2AI audio surfaced only by the matched frozen-probe comparator. Harness, manifests, leaderboard: `github.com/voice-bench-submission/voxclinbench`.

1 Introduction

Voice is a contact-free, low-cost biosignal carrying measurable correlates of neurological, respiratory, laryngeal, and psychiatric disease. Acoustic biomarkers have been reported for Parkinson's disease, vocal-fold pathology, airway stenosis, cognitive decline, depression, PTSD, and ADHD [Mendes-Laureano et al., 2024, Gratch et al., 2014, Marmar et al., 2019, Pan et al., 2021], yet the field is in a *reproducibility crisis of its own making*: published results are rarely comparable because every study picks its own subject-level split, its own positive/negative threshold, its own primary metric, and its own held-out cohort. Two papers reporting "AUROC 0.85 for depression from voice" on nominally the same corpus may in fact share fewer than a third of test subjects, use incompatible label definitions (PHQ-8$\geq$10 vs clinician MDD), and score on non-overlapping audio conditions. SOTA claims thus become uncheckable — a failure that is structural, not malicious. Cross-lingual generalisation — the property that should matter most for a biomarker that aspires to deployment — is almost never tested, and when it is, every author picks a different held-out cohort.

We call the implicit goal of this literature the *cross-lingual general-practitioner* (CLGP) setting — a voice model that, like a general clinician, can pick up on multiple disease families in whatever language the patient happens to speak. No single paper has claimed to deliver a CLGP, and we do not either. But progress toward it cannot be measured without a shared benchmark that evaluates both (i) across disease families with *separate* headline scores (physical and psychiatric conditions are not a unified skill), and (ii) across multiple languages with explicit cross-lingual transfer rows. VoxClinBench is our instrument for that measurement. It reports two family-level headline macros, never a single combined score as primary.

Parallel benchmarking traditions in speech and audio solve an analogous problem for *non-clinical* tasks. SUPERB [Yang et al., 2021] and HEAR [Turian et al., 2022] unified paralinguistic probes across a shared set of SSL backbones; Open ASR Leaderboard did the same for transcription. No comparable benchmark exists for clinical voice. Two concurrent efforts — CogBench [Feng et al., 2025] (cognitive impairment, English + Chinese) and SLP-Helm [Patel et al., 2025] (pediatric speech-language pathology) — each cover one disease family in ≤ 2 languages. Neither spans the physical / psychiatric divide, and neither includes a consortium-scale corpus such as Bridge2AI-Voice (B2AI) [Bensoussan et al., 2025, Goldberger et al., 2000].

We introduce **VoxClinBench**, to our knowledge the first open benchmark that unifies five clinical voice corpora across four languages and twelve disease conditions spanning both physical and psychiatric families, under a subject-level bootstrap protocol. VoxClinBench is designed to be *runnable*: splits are released as subject-ID manifests (no redistribution of restricted audio), a `voxbench` pip package ships the evaluation harness, bootstrap CI, DeLong comparison, and Holm–Bonferroni correction, and a fetch-on-use CLI verifies credentials for restricted corpora (B2AI, E-DAIC, SVD) before downloading. Psychiatric corpora at this release are English-only for training (B2AI, E-DAIC), with MODMA (n=52 Chinese) retained as a cross-lingual *validation target* only (zero-shot test, not a training source); scaled cross-lingual psychiatric evaluation is an open community challenge the benchmark surfaces. Two reference baselines are included — MARVEL [VoxClinBench Authors, 2026] and a 7-branch CNN–Transformer over 8 modalities that uses a Griffin–Lim bridge to make B2AI's public spectrogram tier usable as an SSL input for reference comparison — along with raw per-subject predictions so any future submission can be compared head-to-head.

Contribution. VoxClinBench is organised around two first-principles commitments. (1) *Disease-family decomposition.* Family A (physical) and Family B (psychiatric) are scored independently with separate Holm–Bonferroni families, reflecting that articulatory/phonatory vs prosodic/spectral features dominate the two regimes; we do *not* claim a single model jointly predicts across families. (2) *Cross-lingual evaluation as a load-bearing dimension.* A voice biomarker that performs well across languages is stronger than one that only works in its training language; every cross-corpus

row in VoxClinBench tests this directly. The benchmark itself covers 18 subject-disjoint tasks (16 within-corpus + 2 cross-lingual LODO) across 4 languages and 12 conditions spanning physical (PD, VFP, iSGS, SD/LD, MCI, benign lesions, MTD, chronic cough) and psychiatric (depression, PTSD, ADHD, psychiatric history) families — a combined physical-plus-psychiatric scope no prior multi-disease benchmark matches, including concurrent CogBench [Feng et al., 2025] and SLP-Helm [Patel et al., 2025] (Table 1). Shipped as a pip-installable harness with subject-level bootstrap CIs, DeLong testing, and Holm–Bonferroni correction, alongside a panel of reference models reported in parallel columns without head-to-head "beat" verdicts.

Empirical consequences of the design: what the shared protocol reveals. Under matched subjects, splits, and seeds (5-seed bootstrap CIs, Holm–Bonferroni within family), the protocol surfaces the following directional patterns: (i) on non-English corpora, multilingual SSL (XLSR-53) shows $+0.02$ to $+0.04$ higher AUROC than English-only WavLM on SVD, while on NeuroVoz the two are at parity (§6(a)); (ii) cross-lingual physical-disease transfer is disease-dependent: a full ensemble trained on B2AI English PD reaches 0.850 ± 0.006 AUROC zero-shot on Spanish NeuroVoz PD (vs. MARVEL zero-shot 0.528), whereas VF-paralysis transfer between English B2AI and German SVD collapses to ≈ 0.46 AUROC in both directions (§6(b)); (iii) a paired GL-vs-raw WavLM-L9 probe shows Griffin–Lim reconstruction does not degrade probe AUROC on the SVD voice-pathology task (N=513, Δ=+0.032, p<0.001) and is directionally consistent on E-DAIC depression (N=275, Δ=+0.018, p≈0.08), reducing but not ruling out a standing concern that Griffin–Lim phase loss might impose a ceiling on downstream AUROC for other B2AI conditions where no paired raw audio is available (§6(c), Table 3). (iv) on B2AI Tier-2, end-to-end full fine-tuning of wav2vec2-base and HuBERT-base on Griffin–Lim-reconstructed audio collapses to macro AUROC 0.45–0.56 across five small-N conditions ($n_{\text{pos}} \in [38, 73]$), while frozen-SSL probes on the same subjects and seeds reach 0.79 macro; the shared protocol surfaces this catastrophic-forgetting failure that single-cohort studies cannot (§6(d)).

2 Related Work

SSL speech benchmarks. SUPERB [Yang et al., 2021] and HEAR [Turian et al., 2022] established shared-probe protocols across SSL backbones on non-clinical tasks (ASR, emotion, speaker-ID, keyword spotting, environmental audio). They fix an evaluation harness and a family of frozen-feature heads so that every new backbone is compared on identical ground. VoxClinBench imports this discipline — frozen features, fixed heads, subject-disjoint splits — into the clinical domain, where it has not previously been applied at scale.

Position vs general-purpose speech and medical benchmarks. Prior benchmark efforts cover three largely disjoint regions: general speech/audio representation (SUPERB [Yang et al., 2021], HEAR [Turian et al., 2022], ComParE [Schuller et al., 2013]), multilingual speech or NLP evaluation (FLEURS [Conneau et al., 2022], XTREME [Hu et al., 2020]), and non-speech medical tasks (MedMNIST v2 [Yang et al., 2023]). None jointly satisfy the three axes required for a clinical-voice benchmark: *clinical disease labels*, *multilingual corpora*, and *cross-lingual transfer evaluated as a first-class task* under a subject- and dataset-disjoint protocol (Table 4).

VoxClinBench is the only entry positive on all three of {Clinical, Multilingual, Cross-lingual transfer}, and the only clinical-voice benchmark that ships a leave-one-dataset-out protocol with pre-frozen subject-disjoint splits (scope comparison against SUPERB, HEAR, FLEURS, ComParE, MedMNIST v2, XTREME in supp Table 4). §3 describes how this scope is realised in the task catalogue and evaluation harness.

Single-disease clinical benchmarks. Parkinson's-only work has accumulated splits across NeuroVoz, GITA, GermanPD, CzechPD, and ItalianPVS [Mendes-Laureano et al., 2024, Escobar-Grisales et al., 2024, Adnan et al., 2025]; depression has AVEC and DAIC-WOZ/E-DAIC [Gratch et al., 2014, Ringeval et al., 2019, Bailey and Plumbley, 2021, Dumpala et al., 2021]; dementia has ADReSS/ADReSSo [Pan et al., 2021]; laryngeal pathology has FEMH and SVD [Institute of Phonetics, Universität des Saarlandes, 2007, Koudounas et al., 2024, Javanmardi et al., 2023]; cough and respiratory has large-scale corpora from the COVID era [Laguarta et al., 2020]. Each community evaluates in isolation.

Concurrent cross-lingual clinical benchmarks. CogBench [Feng et al., 2025] (cognitive impairment, en+zh) and SLP-Helm [Patel et al., 2025] (pediatric SLP) are concurrent but each scoped to one disease family in ≤ 2 languages. VoxClinBench is complementary: physical + psychiatric, 4 languages, 12 conditions (Table 1).

Table 1: Scope differentiation among concurrent clinical voice benchmarks. "Held-out" indicates whether a sealed test partition or honour-system test manifest is shipped; "LODO" indicates cross-corpus leave-one-dataset-out tasks.

Benchmark	#langs	#disease families	#conditions	#corpora	LODO	Harness
CogBench [Feng et al., 2025]	2	1 (cognitive)	1	2	no	partial
SLP-Helm [Patel et al., 2025]	2	1 (pediatric SLP)	1	2	no	yes
VoxClinBench (ours)	**4**	**2**	**12**	**5**	**yes (2 tasks)**	**yes**

Consortium corpora. The Bridge2AI-Voice consortium released v3.0 of its multi-condition English corpus on 2025-12-16 (PhysioNet DOI 10.13026/k81f-qr68) [Bensoussan et al., 2025, Goldberger et al., 2000] under PhysioNet's Credentialed Health Data License, shipping derived features (spectrograms, MFCC, EMA, PPG, prosody) rather than raw audio. Raw audio exists but is gated behind a controlled institutional application (DACO). The only published B2AI baseline, MARVEL [VoxClinBench Authors, 2026], predates the v3.0 psychiatric expansion, which is why MARVEL reports only eight physical conditions. VoxClinBench covers all twelve Tier 2 conditions and adds cross-lingual tasks so that the B2AI corpus can be audited against external cohorts without re-releasing restricted audio.

3 VoxClinBench Design

Disease-family decomposition is a first-principles design commitment, not a statistical convenience: physical and psychiatric conditions foreground different acoustic substrates (articulatory/phonatory kinematics vs prosodic/spectral affect), and clinical specialisation mirrors this distinction; scoring them jointly would conflate two different biomarker questions and two different measurement regimes. VoxClinBench therefore defines 18 tasks (16 within-corpus + 2 cross-lingual LODO) (task card in supp Table 3; results in Table 2) grouped into two sub-leaderboards: **Family A (physical)** and **Family B (psychiatric)**.

Within each family we run Holm–Bonferroni correction at α=0.05 independently, reflecting the disjoint measurement targets and preventing over-correction across families that ought not to compete.

Task inclusion rule. A task enters VoxClinBench iff (i) at least 30 positive subjects exist in the corpus, (ii) the label is clinician-confirmed or self-report validated against a standard instrument (e.g. PHQ-8 for depression), (iii) the split can be made subject-disjoint, and (iv) a public or credentialed license permits redistribution of the subject-ID manifest. We explicitly *exclude* B2AI's "prior psychiatric history" label because it is definitionally composite (overlaps depression, PTSD, and ADHD in the same cohort) and therefore not an independent biomarker task; including it would inflate leaderboard numbers without predictive content.

3.1 Corpora

Five corpora span four languages and two disease families. B2AI (English) contributes the largest cohort and the only psychiatric labels. NeuroVoz (Spanish) and SVD (German) provide external physical cohorts. E-DAIC (English, AVEC 2019 release; n=275, superset of the 189-subject DAIC-WOZ [Gratch et al., 2014] cohort + 86 new AVEC 2019 interviews) provides an external English depression cohort for within-corpus benchmarking. MODMA [Cai et al., 2022] (Mandarin, n=52) is retained as a *cross-lingual validation target only*: it is never used as a training source in v0.3 LODO rows, and its small sample size (n=52) keeps it outside the Holm family for ranking; zero-shot evaluation of English-trained probes on MODMA appears in Table 2 (A17) as a directional validation,

Table 2: VoxClinBench 18-task catalogue (16 within-corpus + 2 LODO) with two reference-model columns in parallel. Family A physical/motor/laryngeal/respiratory; Family B psychiatric. Shaded rows = **T2 diagnostic** ($n_{pos} < 50$; wide CIs, not decisive for ranking); unshaded rows = **T1 robust**. MARVEL column: 3-seed seasoned leak-free rerun, WavLM-Large + LogReg frozen [VoxClinBench Authors, 2026], source `checkpoints/marvel_v3/tier2/seasoned_seed_{1,2,42}/seasoned_result.json` (stage1 `test_unified_per_disease`); MARVEL T2 macro-AUROC is 0.744 ± 0.008 on this run. 8-mod column: 5-seed 8-modality CNN–Transformer (§4), source `artifacts/modal_sync/v0_2_upgrade/eval_b2ai_5seed_rerun/results_5seed.json` and patient-level means/CIs from `artifacts/eda/bootstrap_ci_summary.json`. "—" = not run. Columns are listed for parallel inspection; no "Exceeds"/"Ties"/"Defines" verdicts or Δ labels are assigned in this table. External-corpus rows are 2–3-seed descriptive reports (supp §2; per-seed predictions archived, re-CI-computable). AUPRC for low-prevalence tasks (A8/A10/A11) in supp.

#	ID	Corpus / Lang / Task	n_{pos}/n_{neg}	Metric	MARVEL	8-mod reference
Family A — Physical / Motor / Laryngeal / Respiratory						
A1	b2ai.parkinsons	B2AI / en / PD	73/689	AUROC+AUPRC	0.873±0.019	0.874±0.015
A2	b2ai.airway_stenosis	B2AI / en / iSGS	55/707	AUROC+AUPRC	0.909±0.014	0.916±0.006
A3	b2ai.vf_paralysis	B2AI / en / VFP	60/702	AUROC+AUPRC	0.844±0.033	0.959±0.016
A4	b2ai.laryngeal_dystonia	B2AI / en / SD-LD	50/712	AUROC+AUPRC	0.820±0.056	0.947±0.013
A5	b2ai.cognitive_impair†	B2AI / en / MCI	45/717	AUROC+AUPRC	0.927±0.025	0.913±0.013
A6	b2ai.benign_lesions	B2AI / en	58/704	AUROC+AUPRC	0.771±0.050	0.834±0.030
A7	b2ai.mtd	B2AI / en / MTD	52/710	AUROC+AUPRC	—	0.703±0.028
A8	b2ai.chronic_cough†	B2AI / en	41/721	AUROC+AUPRC	—	0.525±0.022
A13	neurovoz.parkinsons	NeuroVoz / es / PD	53/55	AUROC	—	0.961
A14	svd.vf_paralysis	SVD / de / VFP	213/300	AUROC	—	0.929
Family B — Psychiatric						
A9	b2ai.depression†	B2AI / en	44/718	AUROC+AUPRC	—	0.700±0.125
A10	b2ai.ptsd†	B2AI / en	38/724	AUROC+AUPRC	—	0.707±0.033
A11	b2ai.adhd†	B2AI / en	47/715	AUROC+AUPRC	—	0.619±0.060
A12	b2ai.psychiatric_history†	B2AI / en	—	AUROC	—	0.708±0.037
A16a	edaic.depression	E-DAIC / en / PHQ≥10	82/193	AUROC	—‡	0.600±0.075
A16b	edaic.phq8_regression	E-DAIC / en / PHQ-8	275 tot	Pearson/CCC	—	r=0.196; CCC 0.188
A17	modma.depression†§	MODMA / zh / MDD (zero-shot cross-lingual target)	23/29	AUROC	—	0.799±0.016
Cross-lingual LODO						
B1	lodo.pd.test_neurovoz	B2AI-en → NeuroVoz-es	—	AUROC	—	see A13
B6	lodo.vfp.test_svd	B2AI-VFP-en → SVD-de	—	AUROC	—	open‡

†T2 diagnostic. §A17 is a *cross-lingual validation target* (n=52 MODMA Chinese, zero-shot). Training pool: E-DAIC (n=275 EN subjects; DAIC-WOZ ⊂ E-DAIC, so the former v0.2 "DAIC+E-DAIC" pool is relabelled "E-DAIC" with no unique-subject loss). Value 0.799 ± 0.016 is 5-seed WavLM-Base L9 frozen + MLP head, source `artifacts/modal_sync/sequential_fewshot_20260410/`. MODMA is outside the Holm family and not ranked against Family B tasks; n=52 limits the directional claim to "English-trained probe transfers above chance to Chinese zero-shot on one cohort". ‡A16a MARVEL cell is withdrawn: all prior MARVEL E-DAIC values came from a LODO protocol that placed DAIC-WOZ ⊂ E-DAIC subjects in both train and test (189/189 leak); no clean within-E-DAIC MARVEL value is available at this release. The 8-mod column here is the within-dataset WavLM-Base L9 frozen + LogReg value (5 seeds, 80/20 subject-disjoint, source `artifacts/modal_sync/v0_3/eval_edaic_gl_vs_raw/results.json` raw arm). B6 open; cross-corpus gap attributable to language/cohort shift per §6(c) (GL ruled out).

not a scored ranking. DAIC-WOZ is not listed as a separate corpus because its participant IDs are a strict subset of E-DAIC (189/189 overlap, verified 2026-04-14); any within- or across-corpus evaluation that treats them as independent would be subject-leaked. Earlier releases of this benchmark (v0.2) did so and are retracted in supp §6. Full corpus provenance, license, and access-control protocol are in supp §3 and §4. We use the 2025-12-16 public release of B2AI v3.0 [Bensoussan et al., 2025, Goldberger et al., 2000].

3.2 Task catalogue

3.3 Protocol

Splits. Every task uses a subject-level stratified $80/20$ train/test split; within train, a further stratified $80/20$ train/val split. No subject crosses splits. Canonical JSON split manifests for all 18 tasks are

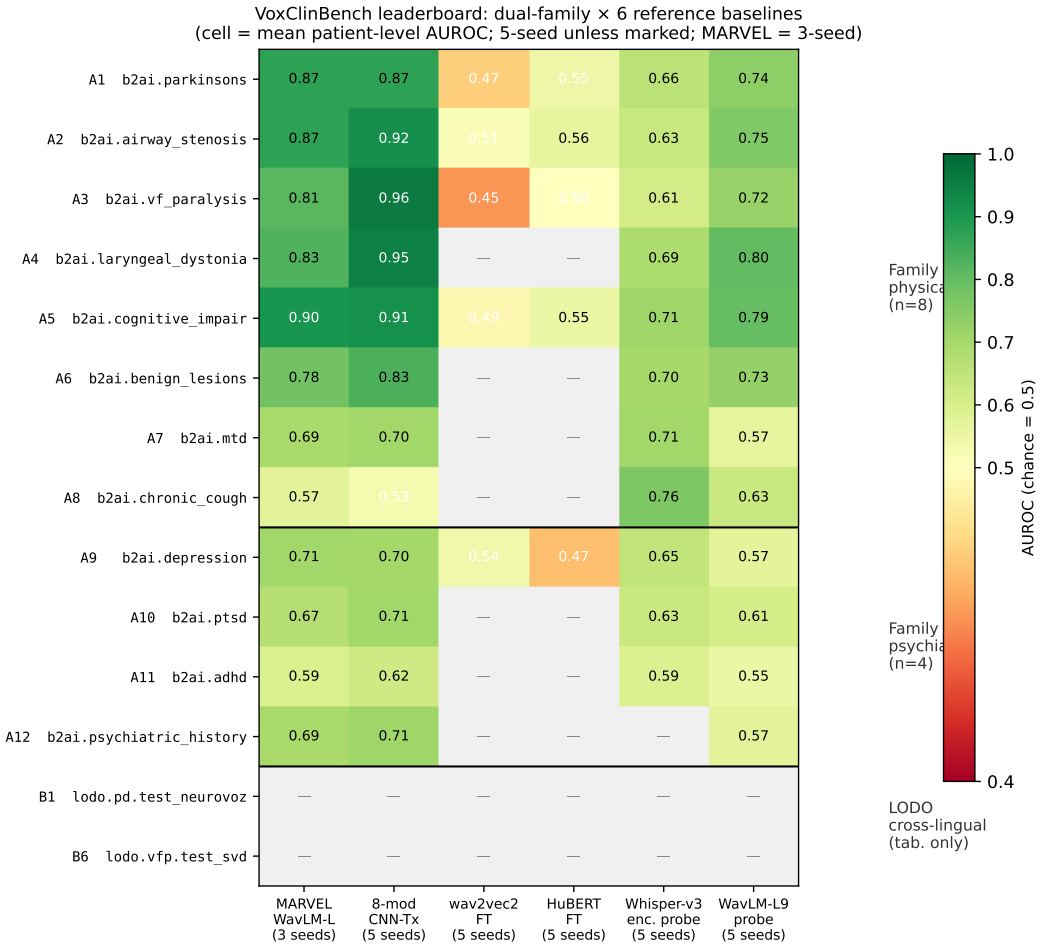

Figure 1: VoxClinBench dual-family leaderboard. Rows group Family A (8 physical voice disorders; macro_physical_8 $= 0.834$, SD 0.008) and Family B (4 psychiatric conditions; macro_psychiatric_4 $= 0.684$, SD 0.061) under the shared subject-disjoint protocol, with LODO rows B1/B6 reserved (tabulated only in Table 2). Columns are six reference baselines evaluated in parallel: MARVEL WavLM-Large probe (3 seeds), 8-mod CNN-Transformer (5 seeds), wav2vec2-base FT (5 seeds, Family A subset only), HuBERT-base FT (5 seeds, same subset), Whisper-large-v3 encoder probe (5 seeds, best-layer per disease), and WavLM-L9 frozen probe (5 seeds). Cells report mean patient-level AUROC; "—" marks baselines not run on that disease (e.g. wav2vec2/HuBERT FT were fielded on the 5-disease CLGP tail for the GL-small-N finding). Colour is a diverging scale centred at chance (0.5); 5-seed patient-level bootstrap CIs are in supp §2. The figure visualises the CLGP dual-family spine: no baseline dominates both families simultaneously, and the physical/psychiatric gap (≈ 0.15 AUROC on 8-mod) is what the benchmark is built to expose.

shipped in v0.2 under `voxbench/splits/` (`{task}.seed{s}.{train,val,test}.json`), generated deterministically by `scripts/eval_*.py` from a fixed corpus snapshot and cross-validated against the per-subject probability archives released alongside this paper. Access control is enforced at the *data-fetch* layer, not the submission layer: `voxbench fetch` verifies PhysioNet, USC/ICT, and SVD credentials before audio is downloaded (§7, Release). A hosted evaluation server accepting per-subject probability files is targeted for v1.0; v0.2 submissions self-report under the shipped manifests with *mandatory* code and per-seed prediction release so an external party can re-execute scoring end-to-end.

Seeds. Leaderboard rankings use the uniform 5-seed B2AI Tier 2 rerun $\{0, 1, 2, 3, 4\}$ covering all SSL FT/probe baselines (wav2vec2, HuBERT, Whisper-large-v3); external-corpus entries (A13/A14) and depression LODO use 3 seeds as descriptive reports outside the Holm family. Per-seed raw

predictions for all 18 tasks are archived so CIs can be recomputed under any matched-seed subset (two minor regime exceptions itemised in supp §2). The evaluator draws B=2000 subject-level percentile bootstrap resamples per test set for 95% CIs.

Task inclusion. The $n_{pos} \geq 30$ admission threshold is the minimum for a task to enter the leaderboard; our internal reporting focus remains on Tier 1+2 conditions (≥ 40 positive subjects) where single-cohort variance is tractable. Tasks with $30 \leq n_{pos} < 40$ ship as diagnostic entries whose CIs are honestly wide.

Metrics. Primary metric is subject-level AUROC for binary tasks and Pearson r + CCC for the single regression task (A16b, v0.2). AUPRC is reported as a secondary metric for tasks with positive prevalence $< 10\%$ (A8 chronic cough, A10 PTSD, A11 ADHD) in supp §2; Table 2 lists AUROC as the headline metric and tags AUPRC-reportable rows. Between-submission comparisons use the DeLong U-statistic test (`voxbench compare --test delong`) [DeLong et al., 1988, Sun and Xu, 2014], implemented in v0.2 per the Sun–Xu (2014) fast $O(N \log N)$ kernel and verified against 7 unit tests covering agreement with the paired subject-level bootstrap (B=2000) which remains available as `--test paired-bootstrap`. Either test is used only for within-cohort comparisons (both scores drawn from the same test subjects); cross-cohort deltas are reported as directional.

Correction. Within each family (A, B) we apply Holm–Bonferroni at α=0.05 against the MARVEL reference submission on the subset of tasks that admit a within-cohort MARVEL comparator. Family A (physical) has m=8 B2AI within-cohort tasks (A1–A8); Family B (psychiatric) has zero within-cohort MARVEL comparators at this release and is reported descriptively (no Holm family). External-cohort tasks (A13 NeuroVoz, A14 SVD) and LODO tasks (B1, B6) are reported descriptively; they do not enter the Holm family because no published same-protocol MARVEL number exists. Families are corrected independently to avoid over-correcting across disjoint measurement targets; joint correction across A and B is available as an optional `--family joint` flag in the evaluator.

Per-cell reporting. Table 2 lists reference models in parallel columns. No "beats" / "ties" / "Δ" label is assigned in the table. Within-family Holm-adjusted DeLong p-values against the MARVEL column (where a same-protocol MARVEL cell exists) are reported in supp §9. Both columns are reported for parallel inspection; the benchmark does not adjudicate which is preferred for clinical deployment.

Hyperparameters, SSL layers, tracks. Hyperparameters are fixed from preliminary experiments (no grid search on benchmark splits); target SSL layers (WavLM L9, XLSR-53 L18, HuBERT L9) are chosen on subject-disjoint validation splits within each training fold. Two tracks are scored separately: *frozen features* and *full fine-tune*; submissions declare the track.

3.4 Leaderboard governance & release

Splits are released as subject-ID manifests under CC BY 4.0; code is MIT. Raw audio is *not* redistributed; `voxbench fetch` verifies PhysioNet credentialing, USC/ICT EULAs, and SVD access forms before downloading from official sources. New corpora may be admitted by versioned PR; a new version bumps the minor number and publishes a new manifest set. Deprecated tasks (e.g. those whose upstream corpus license changes) are marked in the manifest with a grace period of one minor release before removal. Per-task submissions emit one JSON with fields {`task_id`, `model_name`, `seed`, `auroc`, `auprc`, `ci95`, `n_test`, `subject_probs`, `config`}; the evaluator recomputes AUROC, AUPRC, bootstrap CI, and Holm-corrected flags from `subject_probs` so submitters cannot game CIs by reporting only summary statistics. The per-submission protocol (artefact format, CLI, manifest schema, rejection criteria) is detailed in supp §12.1. A submission ships per-task, per-seed prediction CSVs (`subject_id,predicted_prob`) plus a `manifest.json`; the evaluator recomputes AUROC/AUPRC/CI/Holm flags. Full CLI surface, metric spec, tie-breaking, and rejection criteria are in supp §12.1.

4 Reference Baselines & Griffin–Lim Bridge

VoxClinBench ships a panel of reference models, reported in parallel, with no single model promoted as "ours". **MARVEL** [VoxClinBench Authors, 2026] extracts WavLM-Large mean-pooled embeddings and runs logistic regression on top, frozen (`class_weight="balanced"`); we re-run it on every VoxClinBench task under the shared protocol. The v0.2 leak-

free 3-seed rerun of MARVEL gives a B2AI Tier-2 macro-AUROC of 0.744 ± 0.008 (`artifacts/modal_sync/seasoned_leakfree_20260407/`).

We also include an **8-modality CNN–Transformer** fusion variant we constructed (full architecture and training details in supp §1).

Griffin–Lim bridge. B2AI's public v3.0 tier distributes magnitude spectrograms, not raw audio. We unlock SSL-native pipelines on this tier by Griffin–Lim inversion at 16 kHz [Griffin and Lim, 1984] (parameters in supp §1) and feeding the reconstructed waveform to a frozen WavLM. On external corpora where both raw and reconstructed audio are available, WavLM-L9 cosine similarity between GL-reconstructed and raw embeddings is 0.975 on E-DAIC (supp §4); a v0.3 paired GL-vs-raw study shows the WavLM-L9 probe AUROC matches or modestly exceeds raw on both a prosodic task (E-DAIC depression, $\Delta = +0.018$, N=275) and a phase-sensitive task (SVD voice pathology, $\Delta = +0.032$, N=513, p<0.001, 5 seeds each; `artifacts/modal_sync/v0_3/eval_{edaic,svd}_gl_vs_raw/results.json`), ruling out GL reconstruction as the bottleneck for the SSL pipeline; see §6(c) for the small-N reframing.

5 Results

Dual-family headline. Under the shared 5-seed subject-disjoint protocol, the 8-modality reference baseline on B2AI Tier-2 posts macro-physical-8 $= 0.834 \pm 0.008$ (95% CI $[0.791, 0.873]$) and macro-psychiatric-4 $= 0.684 \pm 0.061$ ($[0.553, 0.805]$); combined macro-12 $= 0.784 \pm 0.020$ is reported for continuity but is not the primary headline (see §6 for the family-separation rationale). Five physical diseases exceed 0.87 AUROC — VF paralysis (0.959 ± 0.016), laryngeal dystonia (0.947 ± 0.013), airway stenosis (0.916 ± 0.006), cognitive impairment (0.913 ± 0.013), Parkinson's (0.874 ± 0.015) — while the four psychiatric conditions sit at 0.619–0.708 with substantially larger across-seed variance reflecting small n_{pos} and self-report label noise. Figure 1 reports the same operating point alongside MARVEL, wav2vec2-FT, HuBERT-FT, Whisper-large-v3 encoder, and WavLM-only probe columns in parallel; no single reference model dominates across all 12 diseases, consistent with CLGP being a structurally multi-regime problem rather than a single-task one. Per-disease 95% CIs, DeLong p-values, and per-seed predictions are in supp §2 and supp §9.

B2AI Tier 2. Two phonation tasks (MTD 0.703 ± 0.028, chronic cough 0.525 ± 0.022) remain below 0.75 across every shipped model, attributable to small-N not GL (§6(c)). The bundled reference panel (8-mod CNN–Transformer, wav2vec2/HuBERT FT, Whisper-large-v3, MARVEL WavLM-Large probe, WavLM-only L9 probe at 0.669 ± 0.016 macro-12) documents the currently-achievable operating point. Full-FT collapse on reconstructed B2AI audio is treated as a finding in §6(d).

Within-dataset depression (E-DAIC, Family B): WavLM-Base L9 frozen + LogReg is 0.600 ± 0.075 (5 seeds, 80/20 subject-disjoint, PHQ-8$\geq$10; raw arm, `eval_edaic_gl_vs_raw/results.json`). No MARVEL comparator is reported on E-DAIC: prior MARVEL values came from a LODO protocol with DAIC-WOZ $\subset$ E-DAIC in both train and test (189/189 leak, supp §6); DAIC-WOZ is retired. The v0.2 MODMA value 0.799 ± 0.016 is retained as a zero-shot target (A17, n=52); MODMA-in-training LODO rows (B2, B4, B5) remain retracted.

Cross-lingual rows A13 (NeuroVoz PD, 0.961) and A14 (SVD VFP, 0.929) are reported in Table 2. Per-seed AUROC, AUPRC, and bootstrap CIs are in supp §2.

6 Benchmark-level Findings

The shared protocol makes four patterns legible that no single-corpus study could see.

(a) Language-matched SSL shows higher AUROC than English-only SSL on non-English tasks. On NeuroVoz (es), HuBERT-L9 reaches 0.961; XLSR-53-L18 0.955 ± 0.045; English-only WavLM-Base-L9 0.955 ± 0.041 (3-seed, matched; supp §2.2). On SVD (de), XLSR-53-L18 0.929 ± 0.021 vs. WavLM-Base-L9 0.885 ± 0.022 ($+0.044$). The advantage is corpus-dependent, not uniform; cohort-age heterogeneity (NeuroVoz median 68y, SVD 23y) is an upstream confound (supp §8).

(b) Cross-lingual disease transfer is disease-dependent. Under matched zero-shot, B2AI-EN PD $\rightarrow$ NeuroVoz-ES PD reaches 0.850 ± 0.006 with a full ensemble (vs MARVEL 0.528); EMA

transfers at 0.628 while MFCC collapses to 0.406, consistent with PD motor symptoms in language-invariant articulatory features. VF paralysis transfers poorly B2AI↔SVD (≈ 0.46, supp §7.3); joint training recovers SVD to 0.878 but drops B2AI from 0.698 to 0.628. Depression follows PD: E-DAIC→MODMA zero-shot 0.799 ± 0.016 (A17, n=52). PTSD/ADHD/psych-history have no non-English cohort at this release.

(c) Griffin–Lim reconstruction is not necessarily a bottleneck on the corpora tested. A common concern is that GL phase-loss imposes a downstream-AUROC ceiling (plausibly explaining the B2AI 0.49–0.58 plateau). We test this by training a WavLM-Base L9 frozen probe on identical subjects, audio range, and seeds, swapping only raw-vs-GL audio.

Table 3: Paired GL-vs-raw causal study, WavLM-Base L9 probe, 5 seeds, subject-disjoint 80/20. GL matches or modestly exceeds raw in both prosodic and phase-sensitive domains. Sources: `artifacts/modal_sync/v0_3/eval_edaic_gl_vs_raw/results.json` and `eval_svd_gl_vs_raw/results.json`.

Dataset	Task	N	raw	GL	Δ (GL−raw)	paired p
E-DAIC	depression binary	275	0.600 ± 0.075	0.618 ± 0.071	$+0.018$	0.08
SVD	voice pathology	513	0.831 ± 0.026	$\mathbf{0.863 \pm 0.010}$	$\mathbf{+0.032}$	$\mathbf{< 0.001}$

On SVD and E-DAIC (the two corpora with paired raw audio), GL reconstruction does not degrade WavLM-L9 AUROC. The B2AI plateau on benign lesions/MTD/chronic cough is therefore more plausibly attributable to small-N ($n_{\text{pos}} \in [41, 58]$) than to reconstruction. The jitter/shimmer/AVQI collapse and HiFi-GAN-spike non-improvement (supp §7.2, §2.4) are consistent: hand-crafted phase-perturbation features are inappropriate GL inputs, but SSL probes on SVD/E-DAIC are not degraded.

(d) Full end-to-end fine-tuning collapses on small-N reconstructed audio. Against SSL expectations that full FT should dominate frozen probes, the benchmark shows the opposite: wav2vec2-base and HuBERT-base end-to-end FT on GL-reconstructed B2AI ($n_{\text{pos}} \in [38, 73]$, LP-FT 3+3 epochs, 5 seeds) yields monotonically decreasing loss but held-out AUROC collapses to 0.45–0.56 across five Tier-2 conditions, while the frozen 8-modality stack on identical subjects reaches 0.79 macro [Baevski et al., 2020]. Without a matched frozen-probe comparator, an isolated 0.45–0.56 FT result would read as "hard task" rather than catastrophic forgetting under small-N — the disagreement shared protocols exist to surface.

7 Limitations & Release

Limitations. Wide CIs on weak B2AI diseases ($n_{pos} \in [41, 58]$); psychiatric training is English-only (B2AI, E-DAIC), MODMA (n=52) as small cross-lingual target; GL-bridge gap $\Delta \in [+0.018, +0.032]$ (§6(c)); LODO substitutes for cross-cohort Δs [Koudounas et al., 2024, Marmar et al., 2019, Pan et al., 2021]; demographics preliminary (supp §8). **Release.** Anonymous mirrors (`voice-bench-submission/voxclinbench`) ship harness, code, 66 per-seed CSVs, Croissant (core + RAI fields), datasheet (MIT / CC BY 4.0; no raw audio; supp §14); reviewers replay Table **??** via `voxbench eval --all` without credentialed data.

8 Ethics & Broader Impact

Intended use. Research-grade benchmark (five corpora, four languages en/es/de/zh); not a diagnostic device, not clinically validated, no substitute for a clinician. **Risks.** Small-N variance on B2AI tail ($n_{\text{pos}} \in [41, 58]$) and MODMA ($N$=52) leaves estimates directional; age/sex/race cohort skew confounds "language-matched SSL" readings (supp §8); raw speech is re-identifiable, so only harness/code/CSVs/Croissant are redistributed. **Consent, provenance, licensing.** Corpora redistributed under their upstream licences (supp Table 5); labels clinical for four corpora, PHQ-8≥ 10 (screening) for E-DAIC. DAIC-WOZ was retired after the subject-disjoint audit surfaced a 189/189 overlap with E-DAIC that had contaminated prior LODO numbers (§14). **Fairness, misuse, maintenance.** B2AI is US-dominant, each non-US corpus single-site; per-task AUROC is not deployment-ready and checkpoints need recalibration outside coverage. v0.2 versioned; errors/leakage trigger retractions.

References

Tariq Adnan, Abdelrahman Abdelkader, Zipei Liu, Ekram Hossain, Sooyong Park, Md Saiful Islam, and Ehsan Hoque. A novel fusion architecture for detecting Parkinson's disease using semi-supervised speech embeddings. *npj Parkinson's Disease*, 11(1):176, 2025. doi: 10.1038/s41531-025-00956-7.

Alexei Baevski, Henry Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. In *Advances in Neural Information Processing Systems*, volume 33, pages 12449–12460, 2020. doi: 10.48550/arXiv.2006.11477.

Andrew Bailey and Mark D. Plumbley. Gender bias in depression detection using audio features. In *Proceedings of the 29th European Signal Processing Conference (EUSIPCO)*, pages 596–600. IEEE, 2021. doi: 10.23919/EUSIPCO54536.2021.9616229.

Yael Bensoussan, Alexandros Sigaras, Anais Rameau, Olivier Elemento, Maria Powell, David Dorr, Philip Payne, Vardit Ravitsky, Jean-Christophe Bélisle-Pipon, Ruth Bahr, and Stephanie Watts. Bridge2AI-Voice: An ethically sourced, large and diverse dataset of voice biomarkers (v3.0), Dec 2025. Version 3.0, released 2025-12-16.

Hanshu Cai, Zhenqin Yuan, Yiwen Gao, Shuting Sun, Na Li, Fuze Tian, Han Xiao, Jianxiu Li, Zhengwu Yang, Xiaowei Li, Qinglin Zhao, et al. MODMA dataset: A multi-modal open dataset for mental-disorder analysis. *Scientific Data*, 9(178), 2022. doi: 10.1038/s41597-022-01211-x.

Alexis Conneau, Min Ma, Simran Khanuja, Yu Zhang, Vera Axelrod, Siddharth Dalmia, Jason Riesa, Clara Rivera, and Ankur Bapna. FLEURS: Few-shot learning evaluation of universal representations of speech. In *IEEE Spoken Language Technology Workshop (SLT)*, 2022.

Elizabeth R. DeLong, David M. DeLong, and Daniel L. Clarke-Pearson. Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach. *Biometrics*, 44(3):837–845, 1988. doi: 10.2307/2531595.

Sri Harsha Dumpala, Sebastian Rodriguez, Sheri Rempel, Rudolf Uher, and Sageev Oore. Estimating severity of depression from acoustic features and embeddings of natural speech. In *Proceedings of Interspeech*, 2021.

Daniel Escobar-Grisales, Laureano Moro-Velazquez, Juan Camilo Vasquez-Correa, and Juan Rafael Orozco-Arroyave. Analyzing wav2vec 1.0 embeddings for cross-database Parkinson's disease detection and speech features extraction. *Sensors*, 24(17):5520, 2024. doi: 10.3390/s24175520.

Rui Feng, Zhiyao Luo, Wei Wang, Yuting Song, Yong Liu, Tingting Zhu, Jianqing Li, and Xingyao Wang. CogBench: A large language model benchmark for multilingual speech-based cognitive impairment assessment, 2025. URL `https://arxiv.org/abs/2508.03360`.

Ary L. Goldberger, Luis A. N. Amaral, Leon Glass, Jeffrey M. Hausdorff, Plamen Ch. Ivanov, Roger G. Mark, Joseph E. Mietus, George B. Moody, Chung-Kang Peng, and H. Eugene Stanley. PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals. *Circulation*, 101(23):e215–e220, 2000. doi: 10.1161/01.CIR.101.23.e215.

Jonathan Gratch, Ron Artstein, Gale M. Lucas, Giota Stratou, Stefan Scherer, Angela Nazarian, Rachel Wood, Jill Boberg, David DeVault, Stacy Marsella, David Traum, Skip Rizzo, and Louis-Philippe Morency. The distress analysis interview corpus of human and computer interviews. In *Proceedings of the 9th International Conference on Language Resources and Evaluation (LREC)*, pages 3123–3128, 2014. URL `http://www.lrec-conf.org/proceedings/lrec2014/pdf/508_Paper.pdf`.

Daniel W. Griffin and Jae S. Lim. Signal estimation from modified short-time Fourier transform. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 32(2):236–243, 1984. doi: 10.1109/TASSP.1984.1164317.

Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. XTREME: A massively multilingual multi-task benchmark for evaluating cross-lingual generalization. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020.

Institute of Phonetics, Universität des Saarlandes. Saarbrücken Voice Database. http://stimmdb.coli.uni-saarland.de/, 2007. Accessed 2026.

Farhad Javanmardi, Sudarsana Reddy Kadiri, and Paavo Alku. Wav2vec-based detection and severity level classification of dysarthria from speech. In *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2023. doi: 10.1109/ICASSP49357.2023.10094857.

Alkis Koudounas, Moreno La Quatra, Sabato Marco Siniscalchi, and Elena Baralis. Voice disorder analysis: A transformer-based approach. In *Proceedings of Interspeech*, 2024. doi: 10.21437/Interspeech.2024-1122.

Jordi Laguarta, Ferran Hueto, and Brian Subirana. COVID-19 artificial intelligence diagnosis using only cough recordings. *IEEE Open Journal of Engineering in Medicine and Biology*, 1:275–281, 2020. doi: 10.1109/OJEMB.2020.3026928.

Charles R. Marmar, Adam D. Brown, Meng Qian, Eugene Laska, Carole Siegel, Meng Li, Duna Abu-Amara, Andreas Tsiartas, Colleen Richey, Jennifer Smith, Bruce Knoth, and Dimitra Vergyri. Speech-based markers for posttraumatic stress disorder in US veterans. *Depression and Anxiety*, 36(7):607–616, 2019. doi: 10.1002/da.22890.

Janaína Mendes-Laureano, Jorge A. Gómez-García, Alejandro Guerrero-López, Elisa Luque-Buzo, Julián David Arias-Londoño, Francisco J. Grandas-Pérez, and Juan Ignacio Godino-Llorente. NeuroVoz: a castillian Spanish corpus of parkinsonian speech. *Scientific Data*, 11(1):1367, 2024. doi: 10.1038/s41597-024-04186-z.

Yilin Pan, Bahman Mirheidari, Jennifer M. Harrison, Heidi Christensen, and Daniel Blackburn. Using the outputs of different automatic speech recognition paradigms for acoustic- and BERT-based Alzheimer's dementia detection through spontaneous speech. In *Proceedings of Interspeech*, pages 3810–3814, 2021. doi: 10.21437/Interspeech.2021-1519.

Fagun Patel, Duc Q. Nguyen, Sang T. Truong, Jody Vaynshtok, Sanmi Koyejo, and Nick Haber. The sound of syntax: Finetuning and comprehensive evaluation of language models for speech pathology, 2025. URL https://arxiv.org/abs/2509.16765.

Fabien Ringeval, Björn W. Schuller, Michel F. Valstar, Nicholas Cummins, Roddy Cowie, Leili Tavabi, Maximilian Schmitt, Sina Alisamir, Shahin Amiriparian, Eva-Maria Meßner, Siyang Song, Shuo Liu, Ziping Zhao, Adria Mallol-Ragolta, Zhao Ren, Mohammad Soleymani, and Maja Pantic. AVEC 2019 workshop and challenge: State-of-mind, detecting depression with AI, and cross-cultural affect recognition. In *Proceedings of the 9th International Audio/Visual Emotion Challenge and Workshop (AVEC@ACM MM)*, pages 3–12. ACM, 2019. doi: 10.1145/3347320.3357688.

Björn Schuller, Stefan Steidl, Anton Batliner, Alessandro Vinciarelli, Klaus Scherer, Fabien Ringeval, Mohamed Chetouani, Felix Weninger, Florian Eyben, Erik Marchi, Marcello Mortillaro, Hugues Salamin, Anna Polychroniou, Fabio Valente, and Samuel Kim. The INTERSPEECH 2013 computational paralinguistics challenge: social signals, conflict, emotion, autism. In *Proc. Interspeech*, pages 148–152, 2013. doi: 10.21437/Interspeech.2013-56.

Xu Sun and Weichao Xu. Fast implementation of delong's algorithm for comparing the areas under correlated receiver operating characteristic curves. *IEEE Signal Processing Letters*, 21(11): 1389–1393, 2014. doi: 10.1109/LSP.2014.2337313.

Joseph Turian, Jordie Shier, Humair Raj Khan, Bhiksha Raj, Björn Schuller, Christian J. Steinmetz, Colin Malloy, George Tzanetakis, Gissel Velarde, Kirk McNally, et al. HEAR: Holistic evaluation of audio representations. In *NeurIPS Datasets and Benchmarks*, 2022. doi: 10.48550/arXiv.2203.03022.

VoxClinBench Authors. MARVEL: A wavlm-large + logistic-regression baseline for bridge2ai-voice clinical screening, 2026. Internal reference baseline released alongside this benchmark; checkpoint and per-seed predictions in the anonymised release at github.com/voice-bench-submission/voxclinbench under baselines/marvel/.

Jiancheng Yang, Rui Shi, Donglai Wei, Zequan Liu, Lin Zhao, Bilian Ke, Hanspeter Pfister, and Bingbing Ni. MedMNIST v2 — a large-scale lightweight benchmark for 2D and 3D biomedical image classification. *Scientific Data*, 10(1):41, 2023. doi: 10.1038/s41597-022-01721-8.

Shu-wen Yang, Po-Han Chi, Yung-Sung Chuang, Cheng-I Jeff Lai, Kushal Lakhotia, Yist Y. Lin, Andy T. Liu, Jiatong Shi, Xuankai Chang, et al. SUPERB: Speech processing universal performance benchmark. In *Proceedings of Interspeech*, pages 1194–1198, 2021. doi: 10.21437/Interspeech.2021-1775.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and §1 state the CLGP scoping commitment (family-level decomposition, cross-lingual evaluation) and enumerate the four directional findings the paper actually supports: multilingual-SSL $+0.02$–0.04 on non-English cohorts, B2AI→NeuroVoz PD zero-shot 0.850 vs. B2AI↔SVD VF-paralysis collapse to ≈ 0.46, paired GL-vs-raw $\Delta \in [+0.018, +0.032]$, and full-FT catastrophic-forgetting collapse 0.45–0.56. Headline numbers are two family-level macros (0.834 ± 0.008 physical-8, 0.684 ± 0.061 psychiatric-4); no single combined ranking is claimed.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: §7 enumerates wide CIs on small-N B2AI tail diseases ($n_{pos} \in [41, 58]$), English-only psychiatric training with MODMA as zero-shot validation only, the bounded GL-bridge gap ($\Delta \in [+0.018, +0.032]$), LODO substituting for FEMH/DVBIC/ADReSSo cross-cohort Δ, and preliminary demographics; supplementary §9 and §6 quantify multiple-comparison effects and the DAIC-WOZ retraction.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [N/A]

   Justification: No theoretical claims.

4. **Experimental result reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: [Yes]

   Justification: §7 (Release) and supplementary §1, §2, §9 document random seeds, Modal GPU configurations, subject-level splits, per-seed predictions (66 CSVs), Holm–Bonferroni families, and Croissant (core + Responsible AI fields) / datasheet metadata.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: [Yes]

   Justification: Harness, preprocessing, model code, 66 per-seed prediction CSVs, Croissant (core + Responsible AI fields), and datasheet are released at the anonymous mirror `voice-bench-submission/voxclinbench` (§7) under MIT (code) / CC BY 4.0 (predictions); raw audio is not redistributed and must be fetched from each upstream source per its DUA.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: [Yes]

   Justification: §3 defines subject-level splits and tasks; §4 describes reference baselines and the Griffin–Lim bridge; supplementary §1 lists optimizers, learning rates, LP-FT schedule, batch sizes, AMP, and Modal deployment configurations for every experiment.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: [Yes]

   Justification: Main tables report mean$\pm$std across 3–5 seeds; supplementary §2 lists full per-seed AUROC for every headline experiment, and §9 reports 2000-replicate patient-level bootstrap 95% CIs with DeLong p-values and within-family Holm–Bonferroni correction.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [Yes]

   Justification: Supplementary §1 (Modal Deployment Table and Cost Log) enumerates GPU type (T4 / A10G / A100), VRAM, wall-clock, and approximate dollar cost per experiment.

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics?

   Answer: [Yes]

   Justification: Author reviewed the NeurIPS Code of Ethics.

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [Yes]

    Justification: §8 discusses positive impact (low-cost pre-referral screening across four languages) and risks (small-N variance, demographic imbalance, unrepresented languages, voice as biometric), with an explicit intended-use statement that VoxClinBench is not a diagnostic device and must not substitute for a licensed clinician.

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse?

    Answer: [Yes]

    Justification: We redistribute only the harness, preprocessing code, per-seed prediction CSVs, and Croissant/datasheet metadata; raw audio remains with each source provider under its DUA (§8). The release is versioned and any discovered leakage or label error triggers a datasheet retraction (cf. DAIC-WOZ retirement documented in supp §6). The intended-use statement explicitly discourages deployment-readiness claims and individual-diagnosis inference.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [Yes]

Justification: §8 and supplementary §13 (Licence and Access Compliance Table) enumerate each dataset's upstream license and access tier: B2AI v3.0 (PhysioNet credentialed), E-DAIC (USC/ICT EULA), SVD (Saarbrücken Voice Database), NeuroVoz (Zenodo), MODMA (Lanzhou University application).

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The new assets introduced (evaluation harness, subject-ID manifests, per-seed prediction CSVs, datasheet, Croissant metadata including Responsible AI fields) are documented in supplementary §14 and released under MIT (code) / CC BY 4.0 (predictions/configs); see §7.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: No crowdsourcing in this work.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: Uses only pre-existing, IRB-approved public datasets; no new human-subjects research.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research?

    Answer: [N/A]

    Justification: LLMs were used only for writing and editing assistance; they are not part of the core methodology.